


CS179F: Projects in Operating System

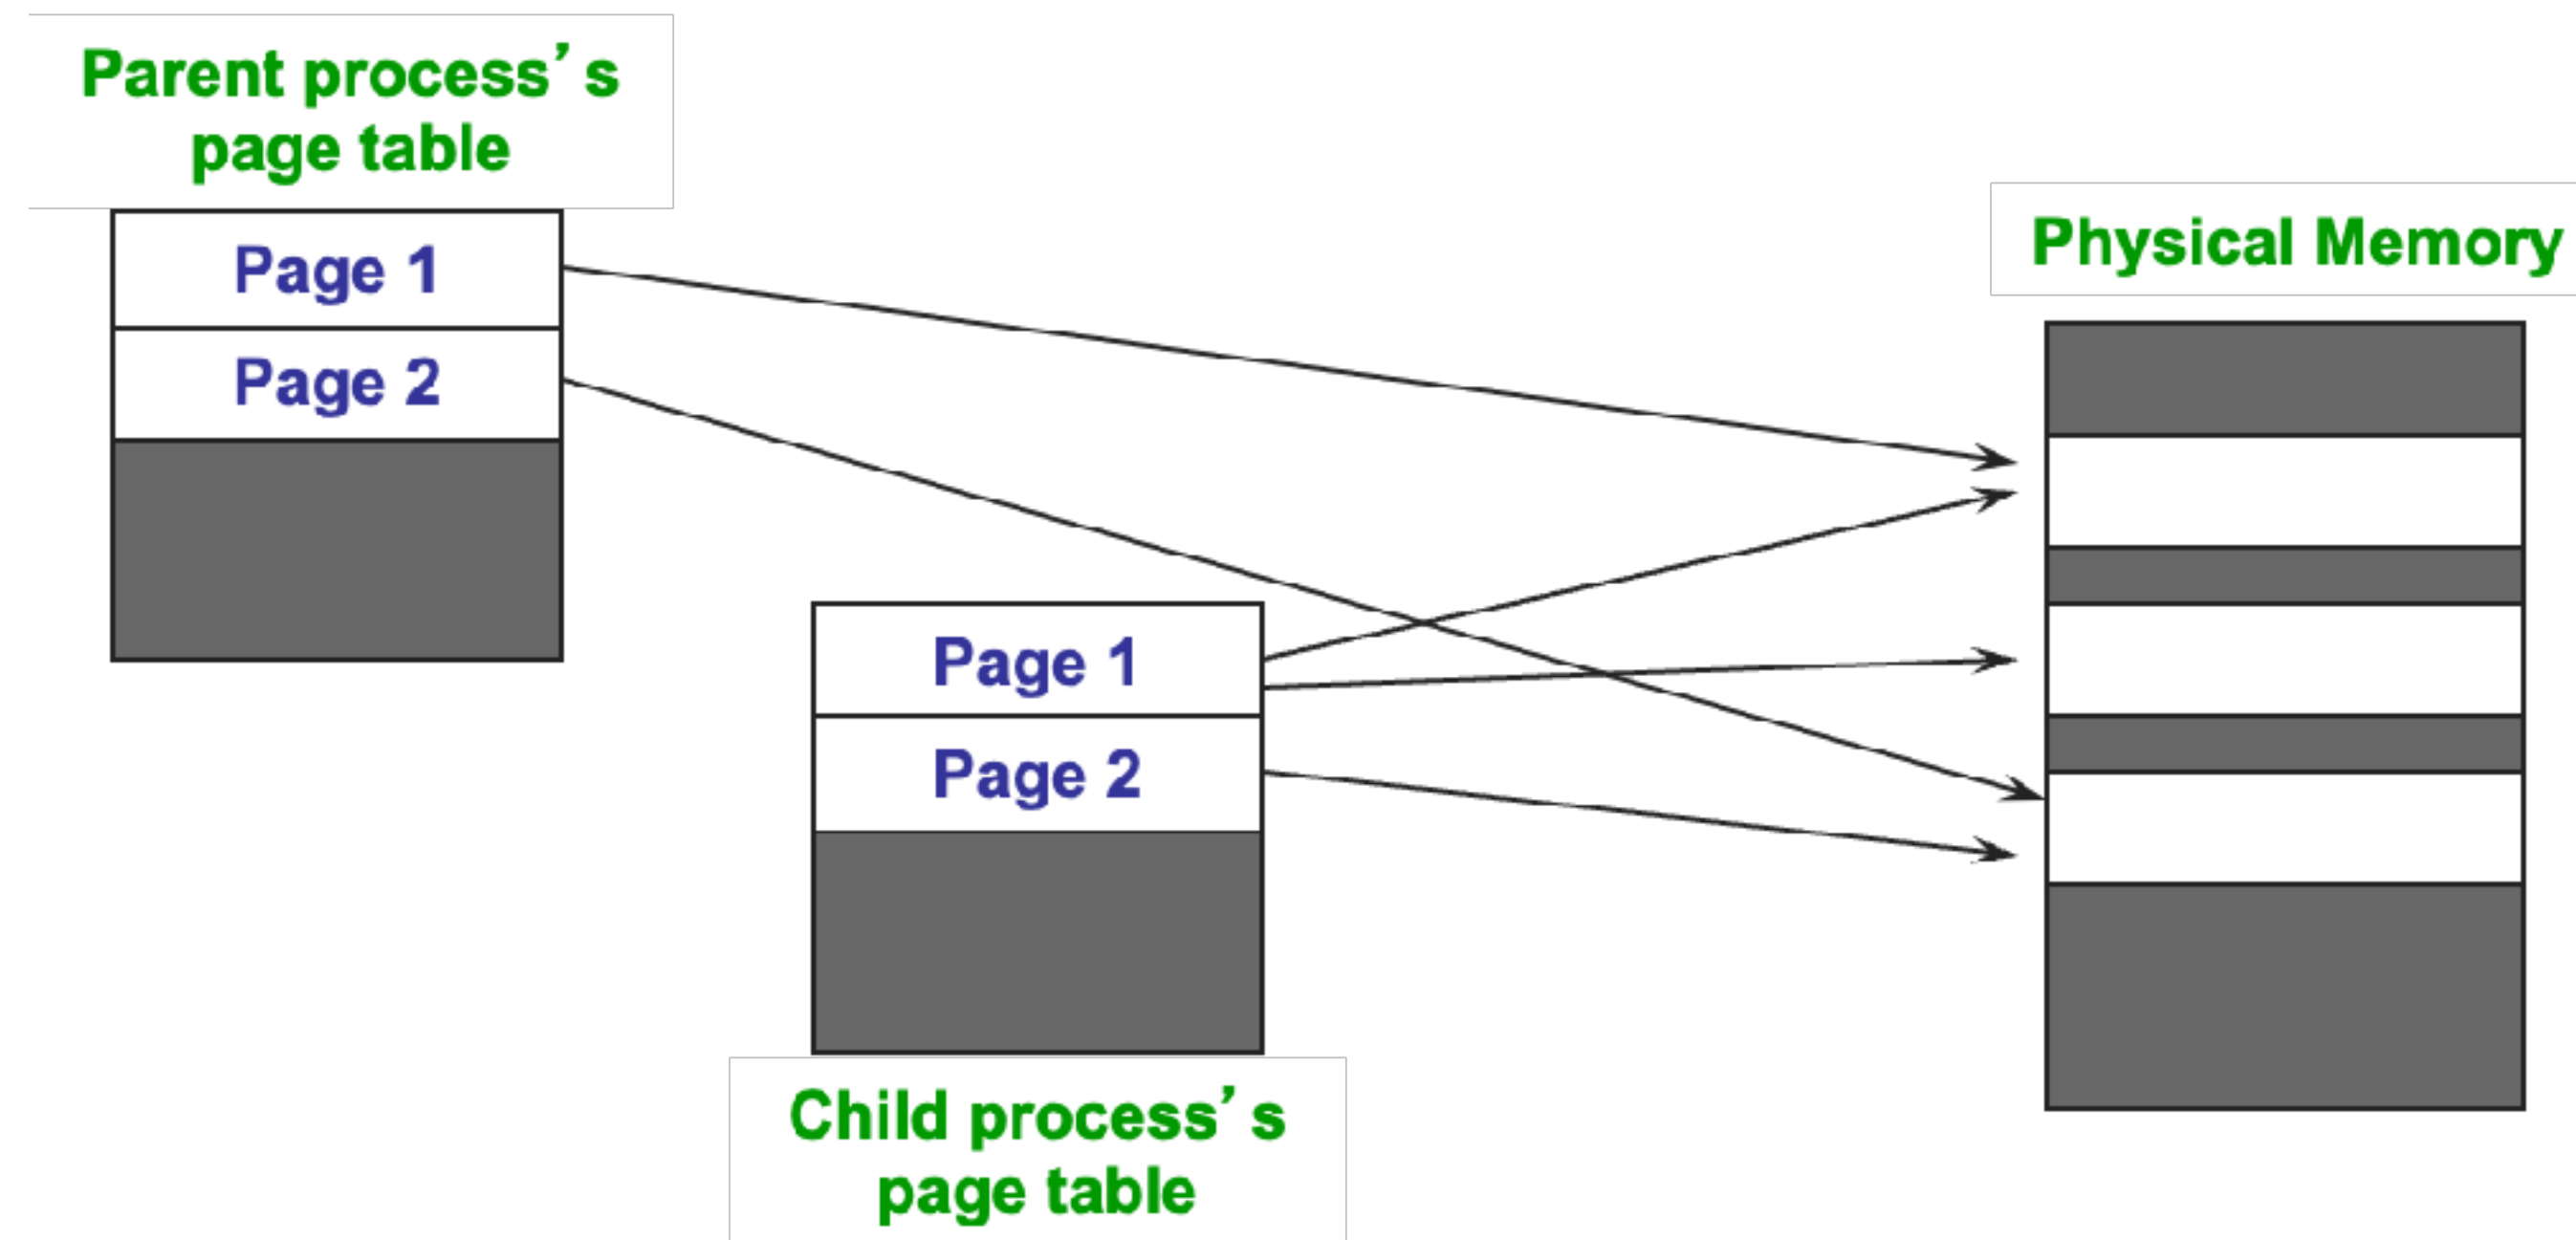
Lab 4: File System

Emiliano De Cristofaro and Lian Gao

fork() without CoW



fork() with CoW



Things to pay attention to

- What is shared?
 - Physical pages, address translations, page tables?
- Page fault handling, how do you know it's caused by CoW?
- When allocating a new physical page, map it to the parent or child's address space?

File system design

- How to allocate and keep track of files and directories?
- Does it matter? What is the difference?
 - Performance, reliability, limitations on files, overhead, ...
- Many different file systems have been proposed and continue to be proposed

Disk layout strategies

- Files span multiple disk blocks
- How do you find all of the blocks for a file?
 - Continues allocation



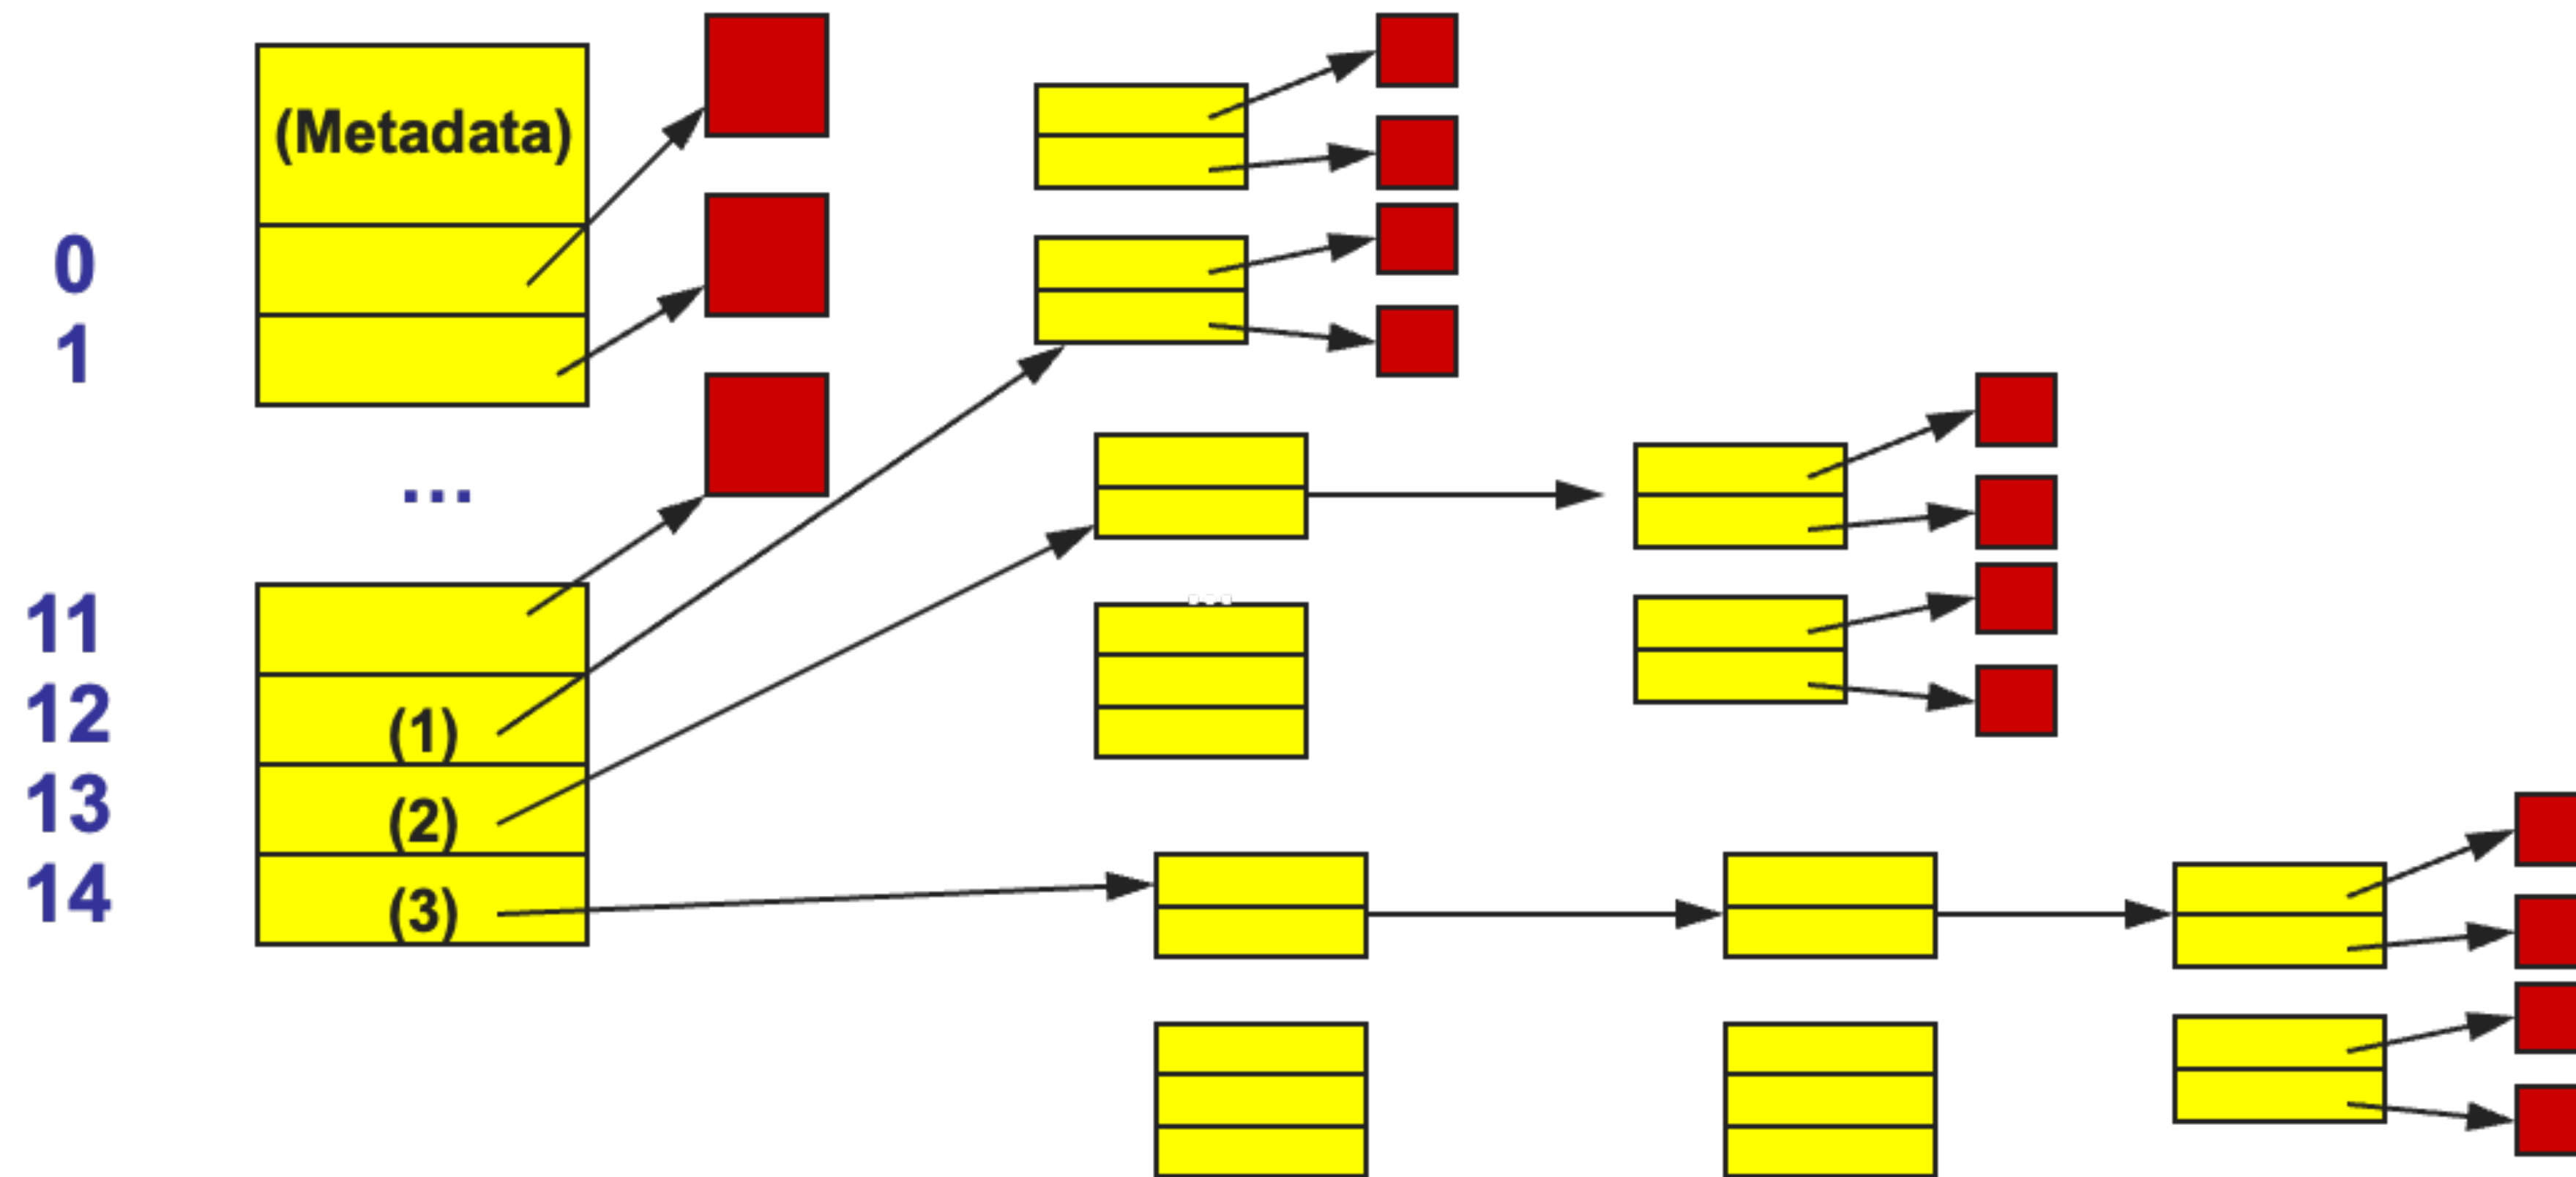
- Linked structure



- Indexed structure (indirection, hierarchy)

UNIX inode

An indexed structure for files



Links

- What is a file?
 - From the OS perspective: the file object (a collection of disk blocks and the corresponding metadata)
 - From users perspective?
- **A link is a pointer to a file**
 - Hard link
 - Symbolic link

Hard Links

Or directory entries

- Hard link is a **reference** to a file object (e.g., an inode)
 - **All named files are hard links!**
 - More than one name can refer to the same file object
- Example
 - `ln myfile mylink`
 - `ls -il`

```
38753 -rw-rw-r-- 2 uli uli 29 Oct 29 08:47 myfile
38753 -rw-rw-r-- 2 uli uli 29 Oct 29 08:47 mylink
^                ^
|-- inode #      |-- link counter (2 links)
```

Hard Links

Operations and Limitations

- How to create/remove a hard link?
 - `link/unlink` <- system calls
- Hard links can only refer to data that exists on the **same** file system
 - Why?
- You **cannot** create hard link to a directory
 - Why?

Symbolic Links

- A **symbolic link** is an indirect pointer to a file – a file whose content is a “pointer/reference” to another “file”
 - “file” actually means a hard link / directory entry
 - so “pointer/reference” means a path
- Example

- `ln -s myfile symlink`

- `ls -li myfile symlink`

```
44418 -rw-rw-r-- 1 uli uli 49 Oct 29 14:33 myfile
44410 lrwxrwxrwx 1 uli uli 6 Oct 29 14:33 symlink -> myfile
^      ^File type ^
|-- inode #      |--link counter (2 links)
```

Symbolic Links

- How does the OS handle symbolic link? (lab4)
 - A symbolic link can be a relative path or an absolute path
- A symbolic link can point to a file on [a different file system](#)
 - Why?
- A symbolic link can point to a non-existent file (referred to as a “broken link”)
 - Why?

Symbolic Links

- You can create a symbolic link to a [directory](#)
 - Why?
- Can we create loops?
 - YES!
 - How to avoid being trapped in a loop?

Lab 4

- xv6 files are limited to 268 blocks, or $268 \times \text{BSIZE}$ bytes (BSIZE is 1024 in xv6)
 - This because an xv6 inode contains 12 "direct" block numbers and one "singly-indirect" block number, which refers to a block that holds up to 256 more block numbers, for a total of $12 + 256 = 268$ blocks.
- In this life, we will learn how to increase the maximum size of an xv6 file
- The bigfile command creates the longest file it can, and reports that size:

Bigfile

- The **bigfile** command creates the longest file it can, and reports that size:

```
$ bigfile
..
wrote 268 blocks
bigfile: file is too small
$
```

- The test fails because bigfile expects to be able to create a file with 65803 blocks, but the unmodified xv6 limits files to 268 blocks.
- Change xv6 to support a "doubly-indirect" block in each in-ode, containing 256 addresses of singly-indirect blocks, each of which can contain up to 256 addresses of data blocks.
- Will create file will be able to consist of up to 65803 blocks, or $256 \times 256 + 256 + 11$ blocks
 - 11 instead of 12 because we will sacrifice one of the direct block numbers for the double-indirect block

mkfs

- Creates the xv6 file system disk image and determines how many total blocks the file system has
- Size controlled by FSSIZE in kernel/param.h (set to 200k blocks in repo)
- You should see the following output from mkfs/mkfs in the make output:

```
nmeta 70 (boot, super, log blocks 30 inode blocks 13, bitmap blocks  
25) blocks 199930 total 200000
```
- This line describes the file system that mkfs/mkfs built: it has 70 meta-data blocks (blocks used to describe the file system) and 199,930 data blocks, totaling 200,000 blocks.
- If you need to rebuild the file system from scratch, you can run `make clean` which forces make to rebuild `fs.img`

What to look at (1/2)

- Format of on-disk inode is defined by struct `dinode` in `fs.h`
 - Look at `NDIRECT`, `NINDIRECT`, `MAXFILE`, and the `addrs[]` element of struct `dinode`
 - Look at Figure 8.3 in the xv6 text for a diagram of the standard xv6 inode
- The code that finds a file's data on disk is in `bmap()` in `fs.c`.
 - Look at it and understand it. `bmap()` is called both when reading and writing a file
 - When writing, `bmap()` allocates new blocks as needed to hold file content, as well as allocating an indirect block if needed to hold block addresses

What to look at (2/2)

- Note: `bmap()` deals with two kinds of block numbers.
 - The `bn` argument is a "logical block number — a block number within the file, relative to the start of the file
 - The block numbers in `ip->addrs[]`, and the argument to `bread()`, are disk block numbers
 - You can view `bmap()` as mapping a file's logical block numbers into disk block numbers

Your job

- Modify `bmap()` so that it implements a doubly-indirect block, in addition to direct blocks and a singly-indirect block
- You'll have to have only 11 direct blocks, rather than 12, to make room for your new doubly-indirect block; you're not allowed to change the size of an on-disk inode
- The first 11 elements of `ip->addrs[]` should be direct blocks; the 12th should be a singly-indirect block (just like the current one); the 13th should be your new doubly-indirect block

Your job

- You are done with this exercise when bigfile writes 65803 blocks and usertests runs successfully:

```
$ bigfile
.....
wrote 65803 blocks
done; ok
$ usertests
...
ALL TESTS PASSED
$
```

- Hints: see <https://github.com/emidec/cs179f-fall23/blob/xv6-riscv-fall23/doc/lab4.md>

Symbolic Links

- Next, you will add symbolic links to xv6
 - Symbolic links (or soft links) refer to a linked file by pathname; when a symbolic link is opened, the kernel follows the link to the referred file
 - Symbolic links resembles hard links, but hard links are restricted to pointing to file on the same disk, while symbolic links can cross disk devices
- Although xv6 doesn't support multiple devices, implementing this system call is a good exercise to understand how pathname lookup works

Your job

- You will implement the `symlink(char *target, char *path)` system call, which creates a new symbolic link at `path` that refers to file named by `target`.
 - For further information, see the man page `symlink`.
 - To test, add `symlinktest` to the Makefile and run it.
- Your solution is complete when the tests produce the following output (including `usertests` succeeding):

```
$ symlinktest
Start: test symlinks
test symlinks: ok
Start: test concurrent symlinks
test concurrent symlinks: ok
$ usertests
...
ALL TESTS PASSED
$
```

Hints: see [git repo](#)