

# CS179F: Projects in Operating System

## Lab 2: Memory Allocation

Emiliano De Cristofaro and Lian Gao

# Dynamic Memory Allocation

- Allocator maintains heap as collection of variable sized *blocks*, which are either *allocated* or *free*
- Types of allocators
  - Explicit allocator: application allocates and frees space (e.g., `malloc` and `free` in C)
  - Implicit allocator: application allocates, but does not free space (e.g., garbage collection in Java, ML, and Lisp)
- We will discuss explicit memory allocation

# malloc-like functions

```
void *malloc(size_t size)
```

- Successful: returns a memory block of **at least** `size` bytes (if `size == 0`, returns `NULL`)
- Failure: returns `NULL`

```
void free(void *p)
```

- Returns the memory block to the pool
- `p` must come from a previous call to `malloc` or `realloc`

# Design Goals


- **Throughput**: number of completed allocation/free requests per unit time
  - We don't want the heap allocator to become the performance bottleneck
- **Memory utilization**: total memory size required to fulfill the requests
  - Fragmentations and poor allocation policies can cause poor memory utilization

# Implementation Issues

- How do we know how much memory to free given just a pointer?
- How do we keep track of the free blocks?
- How do we pick a block to use for allocation -- many might fit?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we reinsert freed block?

# Track Block Size

- **Implicit**: the size can be inferred with other information (e.g., the buddy allocator)
- **Explicit**: record the size information as part of the allocated memory block
  - Requires extra space to store the header



# Keep Tracking of Free Blocks

- **Implicit list:** using a length field to “link” all blocks, and an allocation flag/bit to indicate availability/free



- **Explicit list:** using a (doubled) link list to track all free blocks
- **Separate free list:** an explicit free list for each different size of blocks
- **A heap** (sorted tree): a balanced tree (e.g., red-black tree) to store memory blocks, using size as the key

# Which Block to Pick

- **Frist fit**

- Search the list from the beginning, choose the first free block that fits
- Can take linear time to scan, can cause “splinters” at the beginning

- **Next fit**

- Like first fit, but starts from where the previous search finished

- **Best fit**


- Choose the smallest block that fits (keeps fragmentations small)



# Coalescing

## Defragmentation

- Join (**coalesce**) with next/previous blocks, if they are free



# Freeing With Explicit Free Lists

## Insertion Policy

- LIFO (last-in-first-out)
  - Insert freed block at the beginning of the free list
  - Simple and constant time, but can cause more fragmentations
- Address-ordered
  - Blocks in the free list is ordered by addresses:  $\text{addr}(\text{prev}) < \text{addr}(\text{curr}) < \text{addr}(\text{next})$
  - Requires search, but keeps fragmentations smaller

# The Buddy Allocator

## Split and Merge

Step	34 K	34 K	34 K	64 K	64 K	64 K	64 K	64 K	64 K	34 K	34 K	64 K	64 K	64 K	64 K	34 K
1	$2^1$															
2.1	$2^3$								$2^3$							
2.2	$2^3$				$2^5$				$2^3$							
2.3	$2^1$		$2^5$		$2^5$				$2^3$							
2.4	$2^1$	$2^3$	$2^5$		$2^5$				$2^3$							
2.5	A $2^5$	$2^3$	$2^5$		$2^5$				$2^3$							
3	A $2^5$	$2^3$	B $2^1$		$2^5$				$2^3$							
4	A $2^5$	C $2^5$	B $2^1$		$2^5$				$2^3$							
5.1	A $2^5$	C $2^5$	B $2^1$		$2^1$		$2^1$		$2^3$							
5.2	A $2^5$	C $2^5$	B $2^1$		D $2^1$		$2^1$		$2^3$							
6	A $2^5$	C $2^5$	$2^5$		D $2^1$		$2^1$		$2^3$							
7.1	A $2^5$	C $2^5$	$2^5$		$2^1$		$2^1$		$2^3$							
7.2	A $2^5$	C $2^5$	$2^5$		$2^5$				$2^3$							
8	$2^3$		C $2^5$		$2^5$				$2^3$							
9.1	$2^3$	$2^3$	$2^5$		$2^5$				$2^3$							
9.2	$2^1$		$2^5$		$2^5$				$2^3$							
9.3	$2^3$				$2^5$				$2^3$							
9.4	$2^3$								$2^3$							
9.5	$2^1$															

1. The initial situation.
2. Program A requests memory 34 K.
3. Program B requests memory 66 K.
4. Program C requests memory 35 K.
5. Program D requests memory 67 K.
6. Program B releases its memory.
7. Program D releases its memory.
8. Program A releases its memory.
9. Program C releases its memory.

# A Review on Memory

# Memory

## What is memory


- From programmers' perspective
  - A “place” to store data
- How to access data in memory?
  - Variables?
  - Names?
  - Addresses?
- Memory can be viewed as a big array
  - `content` = `memory[address]`

# Address Spaces

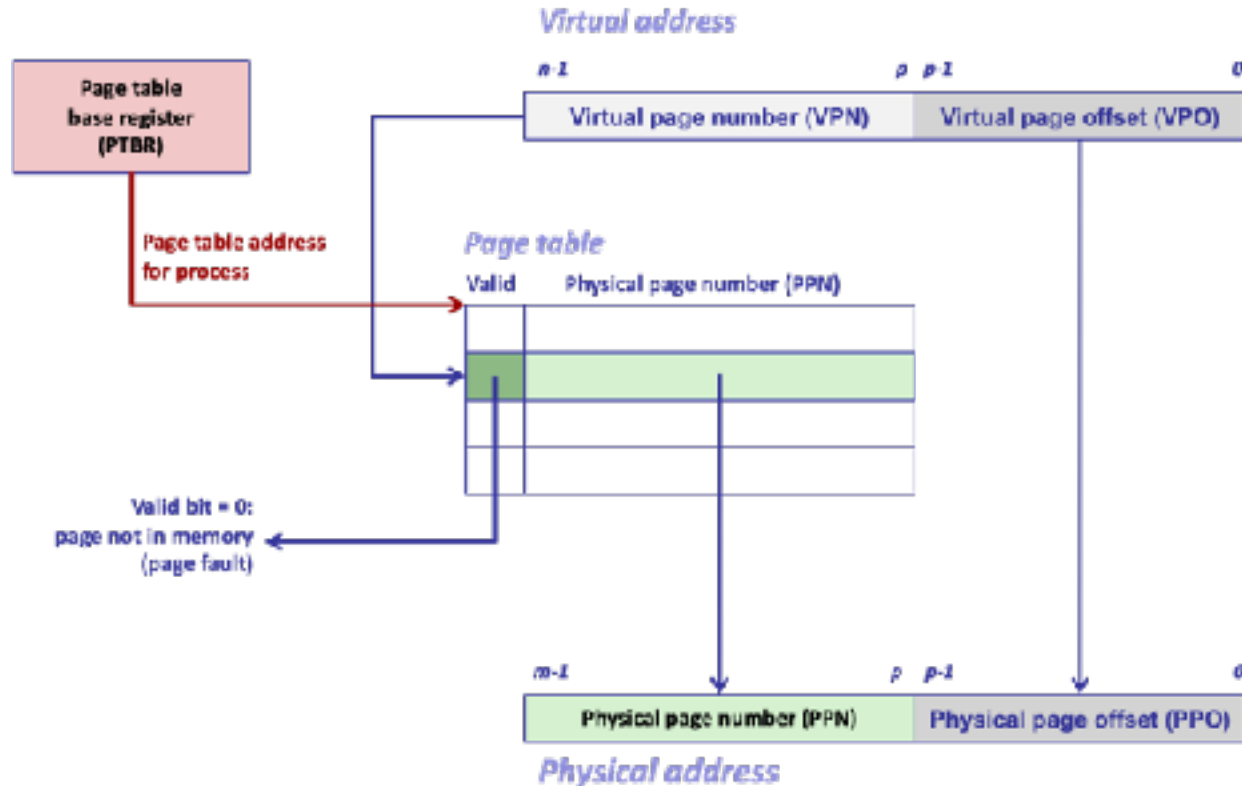
- **Address space**: ordered set of non-negative integer addresses that can be used to access memory: {0, 1, 2, 3 ... }
  - Addresses could be contiguous or segmented
- **Virtual address space**: set of virtual addresses
- **Physical address space**: set of physical addresses

# Paging

- Split the virtual and physical address space into multiple fixed size partitions (i.e., **pages**), each virtual page can be translated to any physical page



# Address Translation with Page Table





# What's wrong in pgbug?

```
int
copyin(pagetable_t pagetable, char *dst, uint64 srcva, uint64 len)
{
    uint64 n, va0, pa0;

    while(len > 0){
        va0 = (uint)PGROUNDDOWN(srcva);
        pa0 = walkaddr(pagetable, va0);
        if(pa0 == 0)
            return -1;
        n = PGSIZE - (srcva - va0);
        if(n > len)
            n = len;
        memmove(dst, (void *)(pa0 + (srcva - va0)), n);

        len -= n;
        dst += n;
        srcva = va0 + PGSIZE;
    }
    return 0;
}
```


```
void
pgbug(char *s)
{
    char *argv[1];
    argv[0] = 0;
    exec((char*)0xeaeb0b5b00002f5e, argv);

    pipe((int*)0xeaeb0b5b00002f5e);

    exit(0);
}
```

# Page Table Entry

## When page fault can happen?



**P:** Child page table present in physical memory (1) or not (0).

**R/W:** Read-only or read-write access permission for all reachable pages.

**U/S:** user or supervisor (kernel) mode access permission for all reachable pages.

**WT:** Write-through or write-back cache policy for the child page table.

**CD:** Caching disabled or enabled for the child page table.

**A:** Reference bit (set by MMU on reads and writes, cleared by software).

**PS:** Page size either 4 KB or 2 MB (defined for Level 1 PTEs only).

**G:** Global page (don't evict from TLB on task switch)

**Page table physical base address:** 40 most significant bits of physical page table address (forces page tables to be 4KB aligned)

**XD:** Non-executable pages

# Page Fault Handling

## OS-induced page faults

- Lazy allocation
  - Aggressive lazy allocation
- Copy-on-Write
- Virtual memory

# Additional Information

- Background:
  - [xv6 book](#): Chapter 2, Chapter 3, and Chapter 4.
  - MIT lecture notes: <https://pdos.csail.mit.edu/6.828/2022/schedule.html>
    - [programming xv6 in C](#)
    - [OS design](#)
    - [Virtual Memory/Page tables](#)
    - [Page faults](#)
- Grading specification:
  - Task 1, pass filetest: **5** points.
  - Task 2, pass lazytests: **10** points.
  - Task 2, pass usertests: **5** points.
- TA will run **ALL** student's git diff, patch and verify the result.

# Grading

- You will get points:
  - Have a **meaningful** report.
    - **screenshots** of the grading script. What tests **passed** and **did not pass**. Your **username**.
    - **explanation** of your code showing your understanding (comments).
  - **Valid** git diff file. **Function level code comments**.
- You will lose points:
  - Do not follow the above instructions.
  - Fail on some tests.
  - Implement something else not asked in the requirement.
- You will **not get any** points:
  - Direct plagiarism.
  - The results reproduced by TA do not match what you present from your report.
  - Invalid git diff file.
  - Invalid submission. Please **upload two files**. report.pdf + mycode.diff